

VR Reading Aid

Nancy Liddle Columbia University
New York, NY, USA
nll12128@columbia.edu




Figure 1: Example use case of the Reading Aid

ABSTRACT

Encountering confusion is an inevitability while learning: if a student comes across text that alludes to or discusses an unfamiliar phrase, algorithm, or event, they often disengage long enough to search for an explanation. This breaks the learning flow and increases cognitive overhead. In order for the brain to retain information, it has to be in a relaxed but attentive state. Immersive environments facilitated by a head-mounted display offer the opportunity to enhance the learning experience and scaffold user cognition. This project presents a Virtual Reality reading assistant that detects moments of elevated mental effort using the HP Reverb G2 Omnicept Edition’s multi-modal sensing and then overlays a context-aware explanation beside the user’s point of fixation. In this abstract, we discuss the implementation of the immersive environment.

Index Terms: Virtual reality, cognitive load, eye tracking, pupillometry, PPG, and large language models

1 INTRODUCTION & RELATED WORK

Virtual Reality has risen in popularity as a potential learning modality due to its ability to facilitate a controlled environment. Prior work shows AR can improve EFL learners’ reading comprehension. [1] Additionally, VR has been reported to modulate brain rhythms associated with learning and memory. [6] Identifying when and where confusion occurred is the first step in enhancing the learning process. Cognitive load is used to describe how much mental effort is required to complete a task. Cognitive Load can infer

memory retention, stress levels, confusion, attention, and learning performance. Prolonged high cognitive load can lead to being overwhelmed and is not conducive to learning. [9] [10] There is extensive literature identifying the predictive powers certain features have for cognitive load, including facial features, pulse, blood pressure, eye movements, pupillometry, EEG signals, and voice volume. [3] Multimodal analytics in VR commonly combines eye tracking and physiological signals to infer user state. [3][8] Recent work detects reading-induced confusion using EEG and eye tracking signals. [11] [5] Wearable sensing in AR has been used to distinguish internal vs. external attention. [4] Researchers have explored the cognitive load classification accuracy of classical machine learning models trained on these features, including: k-nearest neighbor (kNN), naive Bayes (NB), logistic regression, linear discriminant analysis (LDA), support vector machines(SVM), ensemble methods (random forest, XGBoost), and neural networks. For this study, we implemented a VR Reading Aid that builds upon these studies. Contributions. This short paper explores the effectiveness of:

- Scaffolding the learning experience in VR with real-time cognitive load detection.
- Identifying points of confusion by implementing a gaze heatmap over the provided text.
- Gleaning insight from LLM’s after querying for a contextually aware explanation of the identified point of confusion.

2 METHODOLOGY

2.1 Classifying High Cognitive Load

The HP Reverb G2 Omnicept headset has built-in capabilities to estimate cognitive-load levels. The headset receives multimodal

streams of data including pupilometry/eye tracking and pulse (PPG) to classify cognitive load as a continuous value in $[0, 1]$ (with an uncertainty estimate). The cognitive load classifier models weights were trained in a study by Stanford's Virtual Human Interaction Lab. The study had a diverse set of 738 participants complete nine tasks. Each task fell into one of three objective levels of difficulty. Upon the completion of each task, the participant was asked to rate their mental exertion the NASA Task Load Index. Additionally the Participant's performance was scored. For the duration of this study, the participants' eye movements and pulse were recorded. The classifier fuses pupil diameter, position, blink frequency, saccade, and fixation with PPG-derived cardiac features using a CNN-based architecture with attention. This model was found to have a classification accuracy of 78.82% . [7] The weights gleaned from this study were then transferred to the headset, enabling real-time cognitive-load estimates during our study. The user reads text rendered on a virtual book in a controlled VR scene. Should the user's cognitive load exceed a threshold, we identify the most likely source of difficulty on the page using gaze clustering.

2.2 Estimating Region of Confusion with DBSCAN

Let $P = \{p_i\}_{i=1}^N$ be all the gaze points from the previous 30 seconds, and $p_i = (x_i, y_i)$ denote the x, y coordinates of gaze point i.

While the user is reading, we overlay a gaze heatmap onto the view. The gaze heatmap is calculated by the following steps:

1. We first collect all the gaze samples from the previous 30 seconds. Let $P = \{p_i\}_{i=1}^N$ be all the gaze points from the previous 30 seconds, and $p_i = (x_i, y_i)$ denote the x, y coordinates of gaze point i. We must identify where on the text plane each gaze coordinate falls by taking the headset's gaze direction and cast a ray from the camera into the scene. If the gaze intersects with the text plane, we then project that 3D point onto the camera's view to get a normalized 2D gaze location on the screen.
2. We apply the DBSCAN algorithm, to the 2D coordinates, on the normalized gaze coordinates to identify regions where gaze samples are densely concentrated, while treating isolated samples as noise.[2] In DBSCAN, a location is considered a cluster if n neighbors lie within an ϵ radius.

Algorithm 1: DBSCAN clustering of 2D gaze points

Input: $P = \{p_i\}_{i=1}^N$ (last 30s gaze points), $p_i = (x_i, y_i)$; radius ϵ ; minimum neighbors m .

Output: Labels $\ell_i \in \{-1, 1, 2, \dots\}$ for each p_i (noise = -1).

- (a) Initialize $\ell_i \leftarrow 0$ for all i ($0 = \text{unvisited}$). Set cluster id $k \leftarrow 0$.
- (b) For $i = 1$ to N :
 - i. If $\ell_i \neq 0$, continue.
 - ii. Compute the ϵ -neighborhood:

$$\mathcal{N}_i \leftarrow \{j \in \{1, \dots, N\} : \|p_j - p_i\|_2 \leq \epsilon\}.$$

- iii. If $|\mathcal{N}_i| < m$, set $\ell_i \leftarrow -1$ (noise) and continue.
- iv. Otherwise start a new cluster: $k \leftarrow k + 1$, set $\ell_i \leftarrow k$, and set seed set $S \leftarrow \mathcal{N}_i \setminus \{i\}$.
- v. While $S \neq \emptyset$:
 - A. Remove some index q from S .
 - B. If $\ell_q = -1$, set $\ell_q \leftarrow k$ (border point).
 - C. If $\ell_q \neq 0$, continue.
 - D. Set $\ell_q \leftarrow k$.
 - E. Compute $\mathcal{N}_q \leftarrow \{j : \|p_j - p_q\|_2 \leq \epsilon\}$.
 - F. If $|\mathcal{N}_q| \geq m$, update $S \leftarrow S \cup (\mathcal{N}_q \setminus \{q\})$.

-
3. Not all gaze clusters reflect meaningful reading behavior. For each cluster C_k , we compute total dwell time and estimate how many times the gaze returns to that region. Should both metrics surpass our threshold, the cluster is retained.

4. We find the location of confusion by determining the median point within the cluster that had the longest summed duration.

2.3 LLM Query and Rendering

When the HP Reverb G2 Omnicept identifies that cognitive load has surpassed a threshold, we take a screenshot of the user's view (including the overlaid gaze heatmap). We then prompt Open Router the screenshot with the reader's view. The overlaid gaze heatmap allows the prompted LLM to identify where in the text confusion occurred by the intensity of the red. The prompt asks the LLM to first identify the phrase of confusion, explain that phrase, and then explain the phrase within the larger context. The returned explanation is then rendered beside the reader's point of fixation.

3 CONCLUSION & FUTURE STUDY

This work presents a sensor-driven VR reading assistant that detects elevated mental effort, localizes the likely source of confusion using gaze clustering, and overlays context-aware explanations. This project aims to reduce cognitive load to enhance the learning process. Cognitive load is not equivalent to confusion; it can reflect focused effort, novelty, or task engagement. Over-triggering risks annoyance and may distract. LLM outputs can be incomplete or incorrect. However, this project proves the potential of an immersive environment to enhance the learning process by providing aid in real time. The system aims to support learners and multilingual readers. Future research would evaluate learning and productivity gains.

REFERENCES

- [1] S. Ebadi and F. Ashrafabadi. An exploration into the impact of Augmented Reality on EFL learners' reading comprehension. *Education and Information Technologies*, 27:9745–9765, Apr. 2022. doi: 10.1007/s10639-022-11021-8 1
- [2] M. Ester, H.-P. Kriegel, J. Sander, and X. Xu. A density-based algorithm for discovering clusters in large spatial databases with noise. In *Proceedings of KDD-96 (Knowledge Discovery and Data Mining)*, pp. 226–231. AAAI Press, 1996. 2
- [3] M. Kolnes, A. Uusberg, and S. Nieuwenhuis. Broadening of attention dilates the pupil. *Attention, Perception & Psychophysics*, 86(1):146–158, Jan. 2024. doi: 10.3758/s13414-023-02793-3 1
- [4] N. Kosmyna and E. Hauptmann. Are you still watching?: assessing internal and external attention in AR using wearable brain sensing glasses. In N. Argaman, H. Hua, and D. K. Nikolov, eds., *Optical Architectures for Displays and Sensing in Augmented, Virtual, and Mixed Reality (AR, VR, MR) V*, vol. 12913 of *Proceedings of SPIE*, p. 1291310. SPIE, 2024. doi: 10.1117/12.3023316 1
- [5] N. Kosmyna, C.-Y. Hu, Y. Wang, Q. Wu, C. Scheirer, and P. Maes. A pilot study using covert visuospatial attention as an EEG-based brain computer interface to enhance AR interaction. In *Proceedings of the 2021 International Symposium on Wearable Computers (ISWC '21)*, pp. 43–47. Association for Computing Machinery, 2021. doi: 10.1145/3460421.3480420 1
- [6] C. Seydel. Virtual reality boosts brain rhythms crucial for neuroplasticity, learning and memory. UCLA Health, June 2021. Accessed: 2025-12-19. 1
- [7] E. H. Siegel, J. Wei, A. Gomes, M. Oliveira, P. Sundaramoorthy, K. Smathers, M. Vankipuram, S. Ghosh, H. Horii, J. Bailenson, and R. Ballagas. HP omnicept cognitive load database (HPO-CLD) – developing a multimodal inference engine for detecting real-time mental workload in VR. Technical report, HP Labs, 2021. 2

- [8] T. Srirangarajan, P. Wang, and J. N. Bailenson. Multimodal analytics in virtual reality. In J. L. Plass, R. E. Mayer, and G. Makransky, eds., *The Handbook of Learning in Virtual Reality*. MIT Press, 2025. In press. 1
- [9] J. Sweller. Cognitive load during problem solving: Effects on learning. *Cognitive Science*, 12(2):257–285, 1988. doi: 10.1207/s15516709cog1202_4 1
- [10] J. Sweller, J. J. G. van Merriënboer, and F. G. W. C. Paas. Cognitive architecture and instructional design. *Educational Psychology Review*, 10(3):251–296, 1998. doi: 10.1023/A:1022193728205 1
- [11] H. Zhuang, D. Baradari, N. Kosmyna, A. Balyan, C. Albrecht, S. Chen, and P. Maes. Detecting reading-induced confusion using EEG and eye tracking. arXiv, Aug. 2025. doi: 10.48550/arXiv.2508.14442 1